

ECS763P

Natural Language Processing

Week 4

Unsupervised Methods & Latent Variable Models


Matthew Purver
(with material from Jurafsky & Martin)

Unsupervised Methods

- So far our tasks have been **supervised** or **one-class**:
 - Text classification (into known classes)
 - Sequence modelling (with known labels)
 - Language modelling (for one class of “language”)
- Many problems are about discovering classes or distinctions we don’t know in advance:
 - Document clustering
 - Word sense disambiguation
 - Grammar induction
 - Language acquisition
 - Topic modelling
 - Unsupervised tagging (POS, DA, ...)
 - ...

Document Clustering


- Simplest approach is k-means
 - (remember this from Data Mining)



- Algorithm pseudocode:
 - Initialise with k random documents/points as centroids
 - **Expectation** step:
 - For all documents, assign to closest centroid
 - Set of documents sharing a centroid is now a cluster
 - **Maximisation** step:
 - For all clusters, re-calculate centroid
 - Repeat until converged

K-means Clustering

- What do we need?
 - The number of clusters k
 - More advanced algorithms can choose this, see e.g. x-means
 - A suitable distance metric
 - e.g. cosine distance between bag-of-words vectors



$$d_1 = \langle 0.2, 0.8, \dots \rangle$$

$$d_2 = \langle 0.8, 0.1, \dots \rangle$$

$$\cos(d_1, d_2) = \frac{d_1 \cdot d_2}{\|d_1\| \|d_2\|}$$

$$d_2 = \frac{0.2 * 0.8 + 0.8 * 0.1 + \dots}{\sqrt{0.2^2 + 0.8^2 + \dots} * \sqrt{0.8^2 + 0.1^2 + \dots}}$$

- Scaling can be important: normalisation, weights
- e.g. TF-IDF weighting (see Information Retrieval lectures)
 - TF = term frequency = frequency of term in this document
usually smoothed, log-weighted e.g. $tf_t = 1 + \log(f_t)$
 - IDF = inverse document frequency = rarity of term across collection
usually smoothed, log-weighted e.g. $idf_t = \log\left(\frac{N}{1 + df_t}\right)$
- Some luck with initialisation
 - No guarantee of finding a global minimum!
 - Often worth trying multiple different runs


Expectation-Maximisation

- K-means is a special example of a very common general approach
 - **Expectation-Maximisation (EM)**
 - Useful whenever we have two unknowns
 - e.g. unknown labels Q , unknown model parameters Θ
 - This is very often the case! E.g. part-of-speech tagging, parsing, ...
- Alternate steps:
 - E: estimate labels based on current parameters
$$\hat{Q}^t = \operatorname{argmax}_Q F(Q, \hat{\Theta}^t)$$
 - (for k-means: estimate labels given current centroid distances)
 - M: estimate parameters based on current labels
$$\hat{\Theta}^{t+1} = \operatorname{argmax}_{\Theta} F(\hat{Q}^t, \Theta)$$
 - (for k-means: estimate centroid distances given current labels)
- But:
 - No guarantee of finding a global minimum, only local
 - Sensitive to initial conditions, learning parameters

Brown Clustering

- Remember class-based language models:

- words generated from **classes**
- ideally LOCATION, PERSON etc
- what if we don't know the classes?
 - (and/or don't have labelled data)




$$P(w_1 w_2 \dots w_n) = \prod_i P(w_i | c_i) P(c_i | c_{i-1})$$

- Start with V words in V classes
 - E step: merge two classes which give smallest decrease in $P(w_1 \dots w_n)$
 - M step: re-estimate class language model parameters
- Either:
 - Stop when you have desired number of classes C
 - Carry on until all are merged into one big cluster ...

Brown Clustering

- Order of merging clusters gives a tree:



- An added bonus: word vectors!
 - Represent each word by the binary string from root to leaf
 - ‘chairman’ = 0010, ‘president’ = 0011, ‘sprint’ = 101
 - Now we have meaningful word representations
 - Similarity between words \approx vector similarity

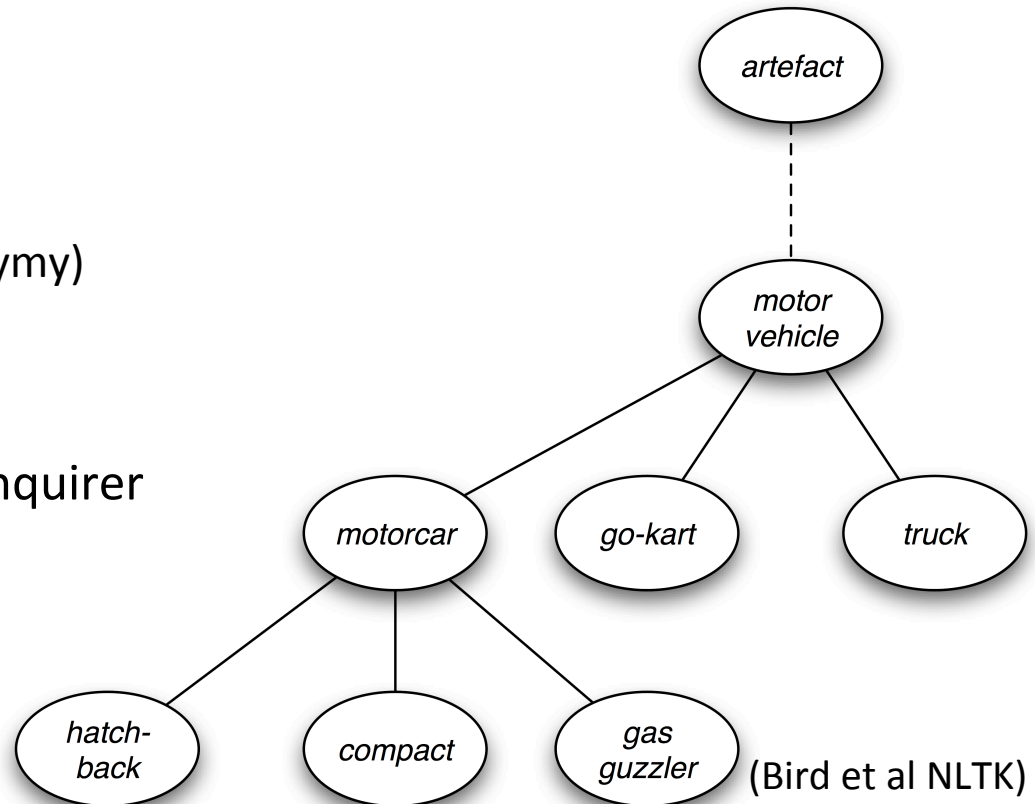
An Aside: Lexical Semantics

- Until now, we've treated words (or n-grams etc) as discrete symbols ...
 - Distinct features in classifier models
 - Distinct symbols in sequence models
- ... and we'll do a lot of this in formal grammars too

mary hires a detective
→ $\exists x. \text{detective}'(x) \ \& \ \text{hire}'(\text{mary}', x)$
- But words themselves have meaning!
 - They can be synonyms, antonyms, hypernyms, hyponyms
good, bad; hot, cold; animal, bird, penguin
 - They can entail each other
bachelor = unmarried, male, human, ...
 - They can be more or less similar to each other
good, bad, terrific, great, awful, terrible
- We are often going to want to know this ...

Lexical Semantics: Ontologies

- Some useful ontology resources exist
 - Good for specific relations
 - Limited, but high-quality
- You might use:
 - WordNet
 - (thesaurus, hypo/hypernymy)
 - VerbNet / FrameNet
 - (verbs & arguments)
 - SentiWordNet / General Enquirer
 - (sentiment & emotions)




Lexical Semantics: Vectors

- Most state-of-the-art NLP now uses vector representations
 - Geometric notions of similarity, negation, entailment etc
 - “distributional semantics” – see later lectures
- Based on language modelling!


- Why does this work?

fill the X, empty the X, ...

- high probability $p(X|w_{1..n}) \approx$ similar meaning




- You might use:
 - Brown clusters (Brown et al 1992)
 - word2vec (Mikolov et al 2013)
 - GloVe (Pennington et al 2014)
 - (pre-trained vectors available for each)



HMM Learning

- We know how to learn HMM models from fully labelled data
- But we usually don't have fully labelled data
 - Why?
- Speech recognition
 - e.g. a phoneme HMM for the word “six”
 - Observations are acoustic features: what do we align it with & how?



HMM Learning

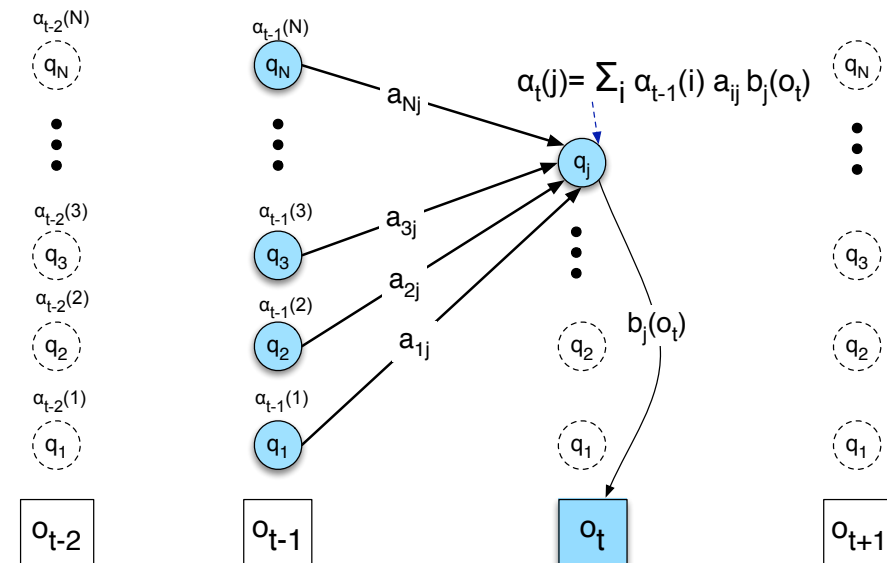
- Speech recognition
 - In fact, we usually don't even know where words start & end:



I want a flight from Boston to Denver

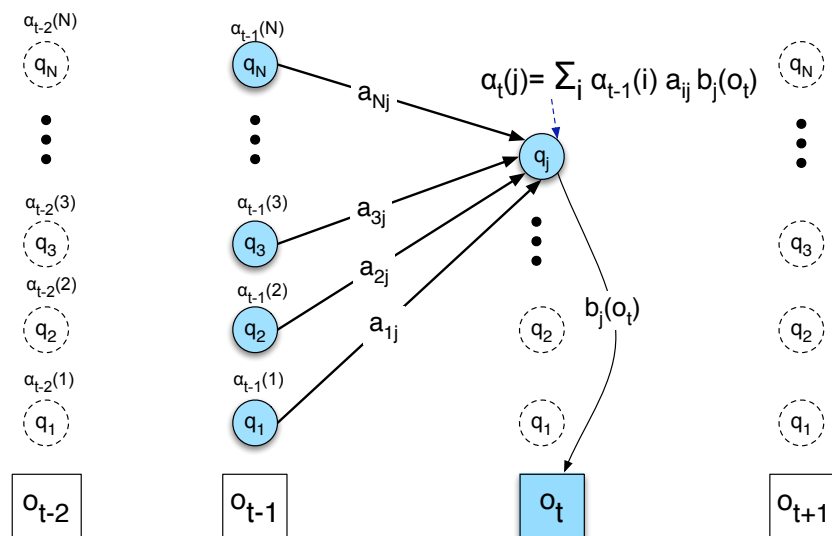
HMM Learning as EM

- So we have two unknowns:
 - State sequence (HMMs are **latent variable** models)
 - Model parameters (transition & emission probabilities A,B)
- Use EM to learn
 - E step:
 - find expected state occupancy and transition counts given A, B
 - M step:
 - re-estimate A, B given estimated counts
- Do this using dynamic programming:
 - Remember Forward algorithm?

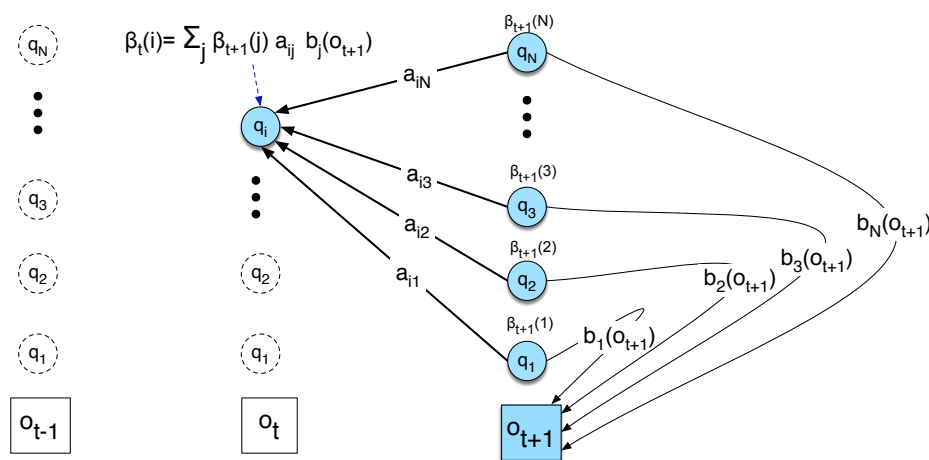


Forward & Backward Probabilities

- “Forward” probability
 - $\alpha_t(j)$ = probability of getting to word t and being in state j (having observed words $o_1 \dots o_t$)



- “Backward” probability
 - $\beta_t(i)$ = probability of getting from word t in state i to the end
 - (observing words $o_{t+1} \dots o_N$)




HMM Learning

- We can combine them in the Forward-Backward (Baum-Welch) algorithm

- Emission probabilities:

- expected number of times seeing word t in state j , from:
$$p(o_t, s_j) = \alpha_t(j) \beta_t(j)$$

- sum over time for each word type, normalise over sum for all word types



HMM Learning


- We can combine them in the Forward-Backward (Baum-Welch) algorithm

- Transition probabilities:

- expected number of times going state i at time t to state j at time $t+1$, from:

$$p(o_{t+1}, s_i \rightarrow s_j) = \alpha_t(i) a_{ij} b_j(o_{t+1}) \beta_{t+1}(j)$$

- sum over time for each state j , normalise over sum for all states



Grammar Induction

- How do we estimate rule probabilities for a grammar?
 - e.g. $S \rightarrow NP VP$ $NP \rightarrow DET N$
 - If we have a manually annotated treebank, we can count: $n(NP \rightarrow DET N)/n(NP)$
 - But of course, usually we don't ...
 - ... and parsing is a very ambiguous process: many possible parses for any sentence and grammar
- Use EM!
 - Initial estimate of rule probabilities
 - E step: parse all sentences, count successful appearances
 - M step: re-estimate rule probabilities
 - Efficient: the Inside-Outside algorithm
- We can even do this to learn the grammar:
 - Invent possible rules
 - E-M to re-estimate rule probabilities, prune unlikely rules

Incremental Grammar Induction

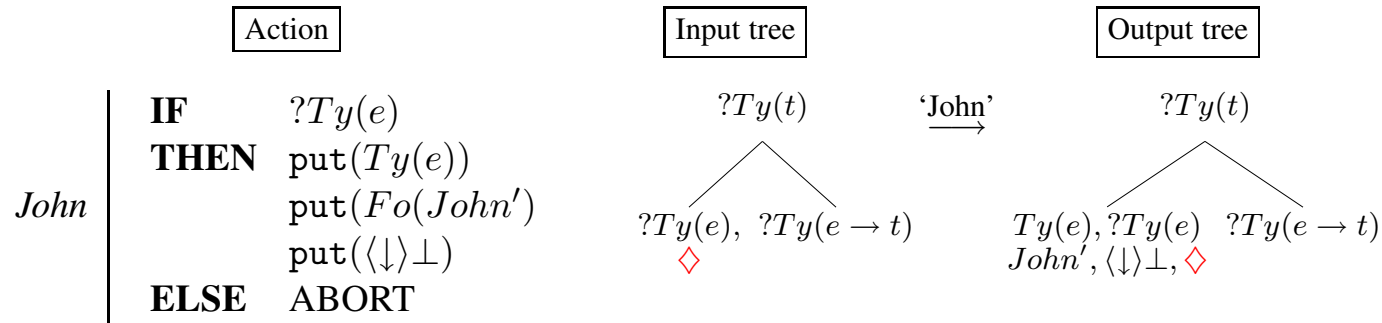


Figure 2: Lexical action for the word 'John'

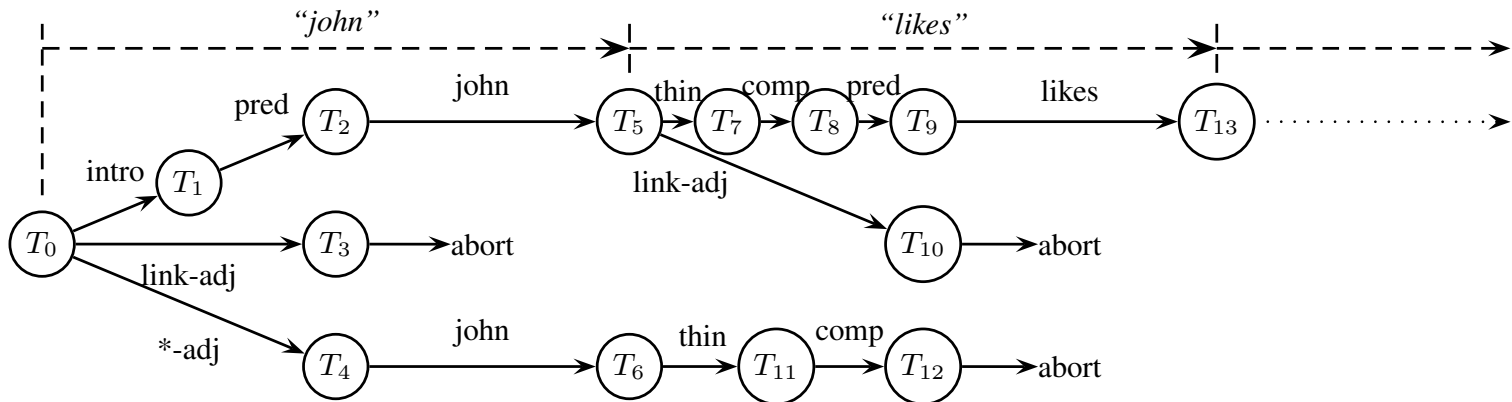


Figure 3: DS parsing as a graph: actions (edges) are transitions between partial trees (nodes).

Topic Modelling

- “What is this document about?”
 - “what is this section about?”
 - “what is this part of this broadcast about?”
 - etc.

[illegible]

- How can we approach this?
 - ~~Most common words~~
 - Most “characteristic” words? (IR)
 - Language model differences?
- But: we’d like something transferable
 - i.e. track topics between documents
- And we’d like some short description
 - i.e. characterise topicality via reduced dimensions

Latent Semantic Analysis

- We already know about dimensionality reduction
 - (remember PCA from Data Mining)
- We can find low-rank approximations to matrices (Deerwester et al, 1988)
 - “matrix factorisation”: usually Singular Value Decomposition (SVD)
- Factorise a term-document matrix X into eigenvectors & eigenvalues :


$$\begin{array}{ccccccc}
 & X & & U & & \Sigma & & V^T \\
 & (\mathbf{d}_j) & & & & & & (\hat{\mathbf{d}}_j) \\
 & \downarrow & & & & & & \downarrow \\
 (\mathbf{t}_i^T) \rightarrow & \begin{bmatrix} x_{1,1} & \dots & x_{1,n} \\ \vdots & \ddots & \vdots \\ x_{m,1} & \dots & x_{m,n} \end{bmatrix} & = & (\hat{\mathbf{t}}_i^T) \rightarrow & \left[\begin{bmatrix} \mathbf{u}_1 \end{bmatrix} \dots \begin{bmatrix} \mathbf{u}_l \end{bmatrix} \right] & \cdot & \begin{bmatrix} \sigma_1 & \dots & 0 \\ \vdots & \ddots & \vdots \\ 0 & \dots & \sigma_l \end{bmatrix} & \cdot & \left[\begin{bmatrix} \mathbf{v}_1 \\ \vdots \\ \mathbf{v}_l \end{bmatrix} \right]
 \end{array}$$

en.wikipedia.org


- Keep the k vectors with highest eigenvalues as “topics”
 - Latent dimensions which explain the variance
- This gives a good model for measuring topical similarity
 - Often more robust than word-based metrics
 - But it’s not very meaningful (how can we understand the “topics”?)

Latent Dirichlet Allocation


- We can build a probabilistic version
 - (but just language modelling isn't enough – why?)
- Generative model (Hoffman, 1999):
 - “probabilistic Latent Semantic Indexing” (PLSI)
 - Assume latent topic variable z
 - Each topic is a word distribution
 - Model document as **mixture** of topics
$$p(w|d) = p(w|z)p(z|d)$$
 - Estimate using EM again
- Gives us interpretable topics!
 - But: no way of applying to a new document
- Latent Dirichlet Allocation (LDA, Blei et al. 2003)
 - Fully generative Bayesian model
 - Assume document-topic mixtures drawn from an underlying distribution
 - (and similarly for topic-word distributions)
 - Estimate using a variant of EM (e.g. Gibbs Sampling)




(a) unigram



(b) mixture of unigrams



(c) pLSI/aspect model



(Blei et al., 2003)

Latent Dirichlet Allocation

- Available LDA implementations e.g. MALLET, gensim

- Sensitive to:

- Number of topics N
- Hyperparameters α, β
 - “spikiness” of distributions
 - How distinct should topics be?
 - How distinct should documents be?
 - Automatic optimisation possible
 - (but be careful!)

- Output:

- N topics as word distributions
- MLE topic distributions for any document


- Often find small number of “junk” topics
 - (see e.g. Griffiths et al 2005 for solutions)

(Howes et al, 2014) – topics from therapy for depression

0	time session sorry today great send next now one work thanks see thank please help make able perhaps look
1	feel life think know way things now like want make self feelings people change maybe someone much need others
2	right well great sure appointment feel thank just lol tonight please know get sorry say bye meeting last though
3	eating eat food weight sick drink meal now lunch control great chocolate absolutely day healthy dinner put use really
4	time husband mum family feel children now dad want see said friends also kids home life got school daughter
5	people say angry situation anger situations said way social others like one friends talk someone person behaviour saying know
6	get go know like need things going just think try want one something time good now make day start

Auto-encoders

- We can achieve similar dimensionality reduction using neural networks:
 - “Autoencoders” – learn optimal reduced encoding from which the input can be reconstructed with some accuracy
 - (more about this later)



Summary

- Estimating models without full supervision:
 - General class of models: latent variable models
 - General method for estimation: Expectation-Maximisation (EM)
- Specific implementations in common areas:
 - Clustering: k-means, Brown etc
 - Generative models: Forward-Backward etc
 - Topic models: pLSI, LDA etc
- No guarantee of exact solution
 - Sensitive to parameters & initialisation
 - Often need multiple runs with
 - Not always obvious how to evaluate & compare